

# Linear Regression Homework

Alireza Afzal Aghaei

March 4, 2021

## Abstract

In this project we have addressed the linear regression model for an specific data-set. The data is collected from a real estate platform in Germany, named Immoscout24. The aim of this exercise is to predict the living space area of a house based on the other features, such as base rent or number of rooms. Here we analyse the data-set features, then remove the ineffective features. The learning process will be done using linear regression model. To increase the model accuracy, we transformed the data into a high-dimensional space using the Nystroem kernel approximation method. The results show a good prediction accuracy for both small and big houses. The mean absolute prediction error for the former is  $3.24m^2$  and  $4.1m^2$  for the latter.

## 1 data-set Description

The data was scraped from Immoscout24, the biggest real estate platform in Germany. Immoscout24 has listings for both rental properties and homes for sale, however, the data only contains offers for rental properties. The scraping process is described in this blog post and the corresponding code for scraping and minimal processing afterwards can be found in this GitHub repository. At a given time, all available offers were scraped from the site and saved. This process was repeated three times, so the data set contains offers from the dates 2018-09-22, 2019-05-10 and 2019-10-08.

The data set contains most of the important properties, such as living area size, the rent, both base rent as well as total rent (if applicable), the location (street and house number, if available, ZIP code and state), type of energy etc. for 268,850 different cases. It also has two variables containing longer free text descriptions: description with a text describing the offer and facilities describing all available facilities, newest renovation etc. The date column was added to give the time of scraping.

In the Table 1, we recall the full column details, provided in the data-set page in Kaggle.

Column Name	Missing	Column Description
regio1	0	Bundesland
serviceCharge	6909	aucilliary costs such as electricity or internet in €
heatingType	44856	Type of heating
telekomTvOffer	32619	Is payed TV included if so which offer
telekomHybridUploadSpeed	223830	how fast is the hybrid inter upload speed
newlyConst	0	is the building newly constructed
balcony	0	does the object have a balcony

**Table 1 continued from previous page**

Column Name	Missing	Column Description
picturecount	0	how many pictures were uploaded to the listing
pricetrend	1832	price trend as calculated by Immoscout
telekomUploadSpeed	33358	how fast is the internet upload speed
totalRent	40517	total rent (usually a sum of base rent, service charge and heating cost)
yearConstructed	57045	construction year
scoutId	0	immoscout Id
noParkSpaces	175798	number of parking spaces
firingTypes	56964	main energy sources, separated by colon
hasKitchen	0	has a kitchen
geo_bln	0	bundesland (state), same as regio1
cellar	0	has a cellar
yearConstructedRange	57045	binned construction year, 1 to 9
baseRent	0	base rent without electricity and heating
houseNumber	71018	house number
livingSpace	0	living space in sqm
geo_krs	0	district, above ZIP code
condition	68489	condition of the flat
interiorQual	112665	interior quality
petsAllowed	114573	are pets allowed, can be yes, no or negotiable
street	0	street name
streetPlain	71013	street name (plain, different formatting)
lift	0	is elevator available
baseRentRange	0	binned base rent, 1 to 9
typeOfFlat	36614	type of flat
geo_plz	0	ZIP code
noRooms	0	number of rooms
thermalChar	106506	energy need in kWh/(m <sup>2</sup> a), defines the energy efficiency class
floor	51309	which floor is the flat on
numberOfFloors	97732	number of floors in the building
noRoomsRange	0	binned number of rooms, 1 to 5
garden	0	has a garden
livingSpaceRange	0	binned living space, 1 to 7
regio2	0	District or Kreis, same as geo krs

Table 1 continued from previous page

Column Name	Missing	Column Description
regio3	0	City/town
description	19747	free text description of the object
facilities	52924	free text description about available facilities
heatingCosts	183332	monthly heating costs in €
energyEfficiencyClass	191063	energy efficiency class (based on binned thermalChar, deprecated since Feb 2020)
lastRefurbish	188139	year of last renovation
electricityBasePrice	222004	monthly base price for electricity in € (deprecated since Feb 2020)
electricityKwhPrice	222004	electricity price per kwh (deprecated since Feb 2020)
date	0	time of scraping

Table 1: Column descriptions

As you can see, some of the scraped features are completely unrelated to our regression task. For example, `scoutId`, `condition`, `street`, `facilities`, `description` and `date` are among the redundant columns.

The next important point is the missing values of the data-set. The Table 1 also report the count of the number of missing values for each feature. There are different methods to handle the missing values in the data-set. dropping rows or columns, imputing them with a constant value or the most frequent value or even filling them with the most similar cases having the desired column. In this exercise we use a more efficient method based on the data-set properties. Before doing it, we should done some preprocessing steps on the data-set. Here we show the distribution of the numeric variables using a box plot. Figure 1 shows this box plot for `floor`, `noRooms`, `livingSpace`, `heatingCosts`, `totalRent` and `baseRent`. You can see there are many outliers in the data. Removing this cases from the data reduces the data-set size to approximately 255000 rows.

The next important information which can help the regression model, is the **livingSpaceRange** feature. This feature summarizes the `livingSpace` feature into 7 different categories. The first category contains the smallest houses while the category 7 includes the biggest living space houses.

From now we preprocess the data based on this feature. Here we explain the outlier detection for the first living space range. The process is same for other ranges, except the bound constraints. Since there are many missing values in columns, we fill them based on their living space range. In other words, we fill the missing values in `heatingCosts` feature, using the values corresponding to the first living space range. This will help filling them with a more appropriate values.

As described in 1, the `totalRent` feature is sum of the base rent, heating cost and service charge. Also,

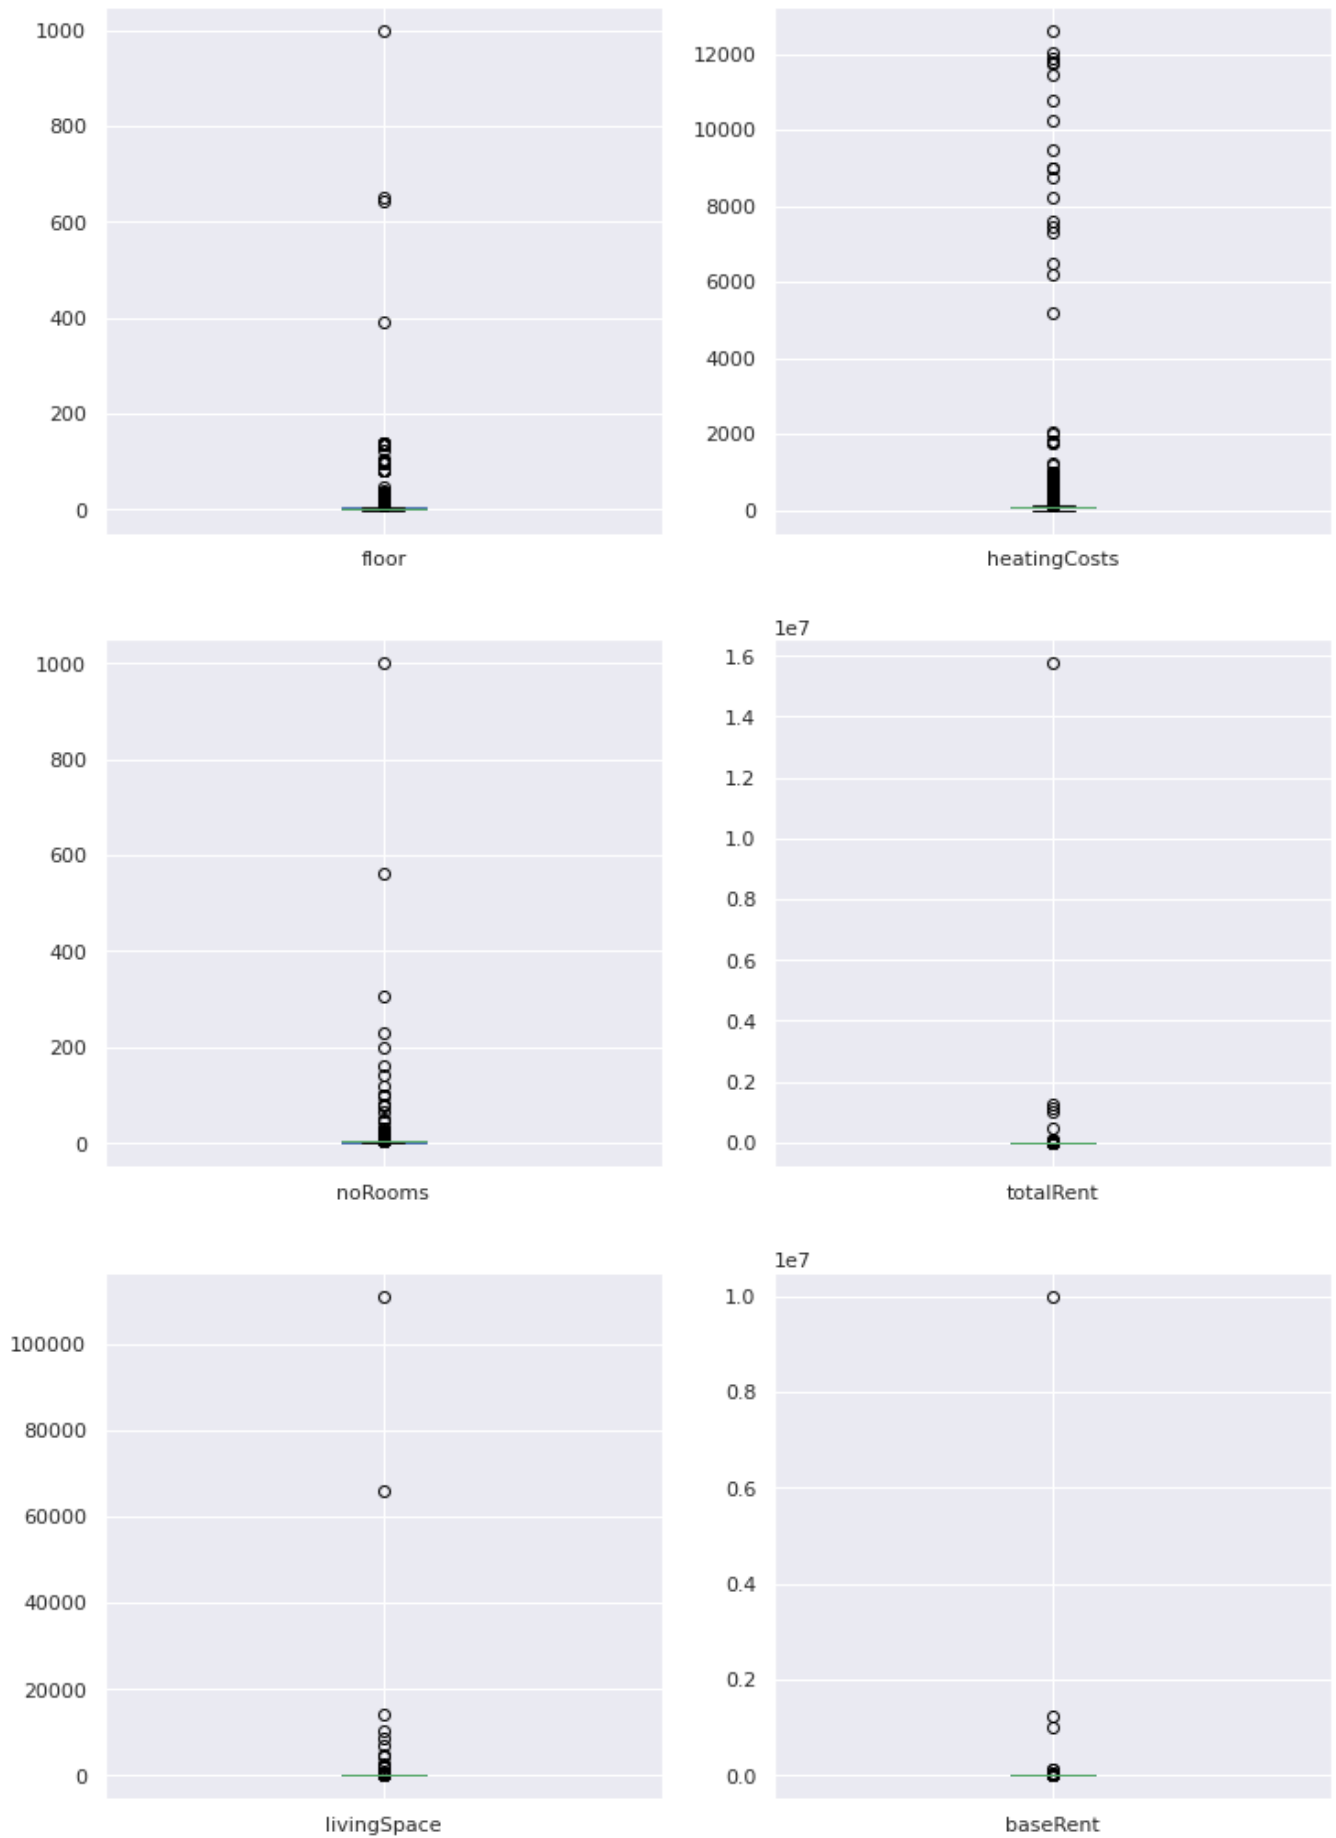


Figure 1: box plot for some of the data-set features

regio1		heatingType		petsAllowed	
Category	Count	Category	Count	Category	Count
Nordrhein_Westfalen	31480	central_heating	77802	negotiable	94758
Sachsen	28490	district_heating	13262	no	20242
Sachsen_Anhalt	11583	gas_heating	9498	yes	5080
Niedersachsen	7188	self_contained_central_heating	7900		
Bayern	6229	floor_heating	5656		
Hessen	5138	oil_heating	2450		
Thüringen	4883	heat_pump	1064		
Baden_Württemberg	4669	combined_heat_and_power_plant	838		
Brandenburg	4107	night_storage_heater	678		
Mecklenburg_Vorpommern	4068	wood_pellet_heating	399		
Schleswig_Holstein	3506	electric_heating	336		
Rheinland_Pfalz	3278	stove_heating	125		
Berlin	2672	solar_heating	72		
Bremen	1336				
Hamburg	840				
Saarland	613				

Table 2: The number of rows in each category

we fill the missing values in **totalRent** feature by the summation of this three feature, based on the living space range.

Now, it's time to remove outliers in this living space range. Again, we draw the box plot for some of the numerical features. Figure 2 shows these plots.

Applying the outlier removal constraints, discards approximately 15000 rows of data. We do the same process for other ranges 2-7. Our analysis showed that the last house range are very difficult to approximation, i.e. they are luxury houses, so we removed the last range. After doing the purity process, we lost about 140000 rows.

The next preprocessing step is to remove the less frequent categories in categorical features. To find them we count unique values in **regio1**, **heatingType**, **petsAllowed** and reported in 2.

Here, we remove 3 less often categories of **regio1**. For the **heatingType** feature we only use the 5 most used categories. Negotiable option in **petsAllowed** feature is mapped to yes, while the missing values considered as no.

For the final preprocessing steps, we round the target feature living space with one digits of accuracy and encode the categorical features using one hot encoding. The data is now ready to be used for our

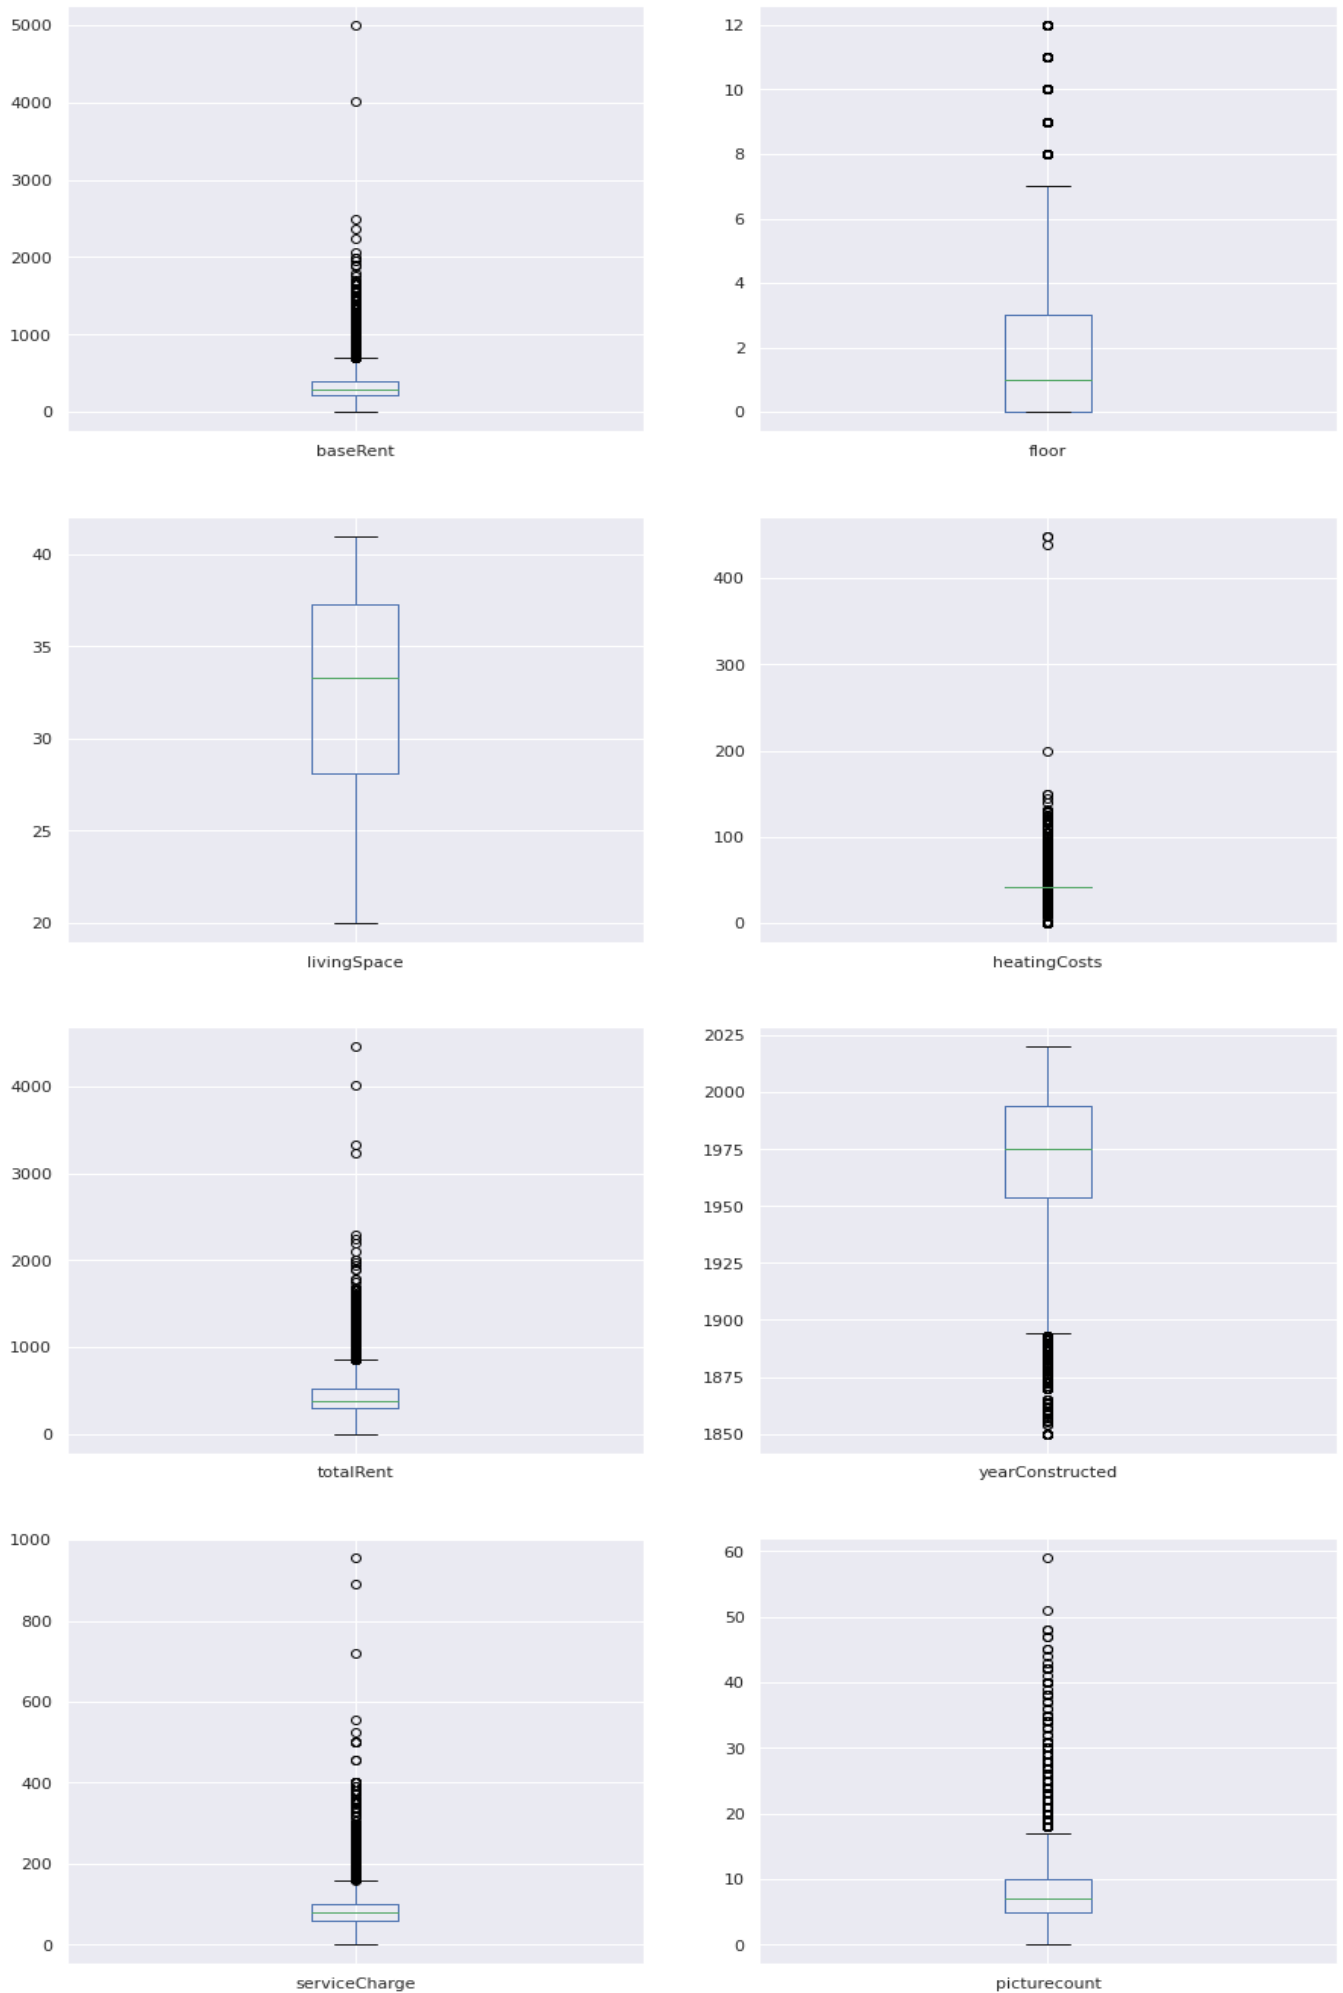


Figure 2: box plot for some of the data-set features

regression task. In the next section we analyse the features using some statistical tests.

## 2 Data analysis

In the previous section we done some preprocessing tasks to remove outliers from the given data-set. In this section, we analyse the final data-set by some statistical measures and tests.

### 2.1 Correlation analysis

For the first test, we compute the correlation matrix of the numerical features of the data-set. The features are `livingSpaceRange`, `heatingCosts`, `totalRent`, `baseRentRange`, `baseRent`, `noRooms`, `noRoomsRange`, `serviceCharge`, `picturecount`, `balcony`, `regio1`, `heatingType`, `petsAllowed` and `livingSpace`. The figure 3 shows this matrix.

As you can see in the plot, the features `totalRent`, `baseRentRange` and `baseRent` are highly correlated. This is not a surprising fact because the base rent range is computed directly based on base rent and the total rent always includes the base rent features. `noRooms` and `noRoomsRange` are highly correlated which is obvious. Finally, the living space and the living space range feature have a high correlation which is directly consequence of categorizing it to 7 bins.

### 2.2 Hypothesis testing

Here we use the ANOVA test to checking whether the two features have the same distribution or not. The ANOVA, tests whether the means of two or more independent samples are significantly different. It has the following assumptions on the data

- Observations in each sample are independent and identically distributed (iid).
- Observations in each sample are normally distributed.
- Observations in each sample have the same variance.

The interpretation of the method is as following:

- H0: the means of the samples are equal.
- H1: one or more of the means of the samples are unequal.

In the table 3 we reported this test on all combinations of the columns of length 2.

## 3 Building the model

The purpose of the exercise is to build a linear regression model to predict the living space area. Given the training data  $X$  and  $y$ , this model defines it's predictor function as

$$y(x) = w^T x + b, \tag{1}$$




Figure 3: box plot for some of the data-set features



Feature 1	Feature 2	ANOVA test
livingSpaceRange	noRooms	+
livingSpaceRange	totalRent	+
livingSpaceRange	serviceCharge	+
livingSpaceRange	livingSpace	+
livingSpaceRange	heatingCosts	+
livingSpaceRange	balcony	-
livingSpaceRange	baseRent	+
noRooms	totalRent	+
noRooms	serviceCharge	+
noRooms	livingSpace	+
noRooms	heatingCosts	+
noRooms	balcony	+
noRooms	baseRent	-
totalRent	serviceCharge	+
totalRent	livingSpace	+
totalRent	heatingCosts	+
totalRent	balcony	+
totalRent	baseRent	+
serviceCharge	livingSpace	+
serviceCharge	heatingCosts	+
serviceCharge	balcony	+
serviceCharge	baseRent	+
livingSpace	heatingCosts	+
livingSpace	balcony	+
livingSpace	baseRent	+
heatingCosts	balcony	+
heatingCosts	baseRent	+
balcony	baseRent	+

Table 3: ANOVA test results. plus means the columns have same distributions, minus means different distributions

where  $w$  and  $b$  are the unknown coefficients. If we define the loss function as

$$loss(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2, \quad (2)$$

the solution of the model can be found by solving the following linear system of equations:

$$\begin{bmatrix} \mathbf{1} & X \end{bmatrix} \begin{bmatrix} b \\ w \end{bmatrix} = y \quad (3)$$

## 4 Fitting the model

Before fitting the model, we split the data-set into two categories: small and big houses. Small cases are those that their living space range are 1 or 2. every house with higher living space range will be considered as a big house. This process helps us to find a more accurate model for each case. It's worth to note that we can use 6 different models for each living space range but it has more computational cost which may not be efficient in some cases.

Now, it's time to fit the model on these sub-data-sets. Before doing this, we transform the raw data into a high dimensional space using the Nystroem kernel approximation method. This method approximates a kernel map using a subset of the data. Here we used a radial basis function as the kernel of the method. The RBFs hyper-parameter  $\gamma$  for both model is found by trial and test. Using this method, we create 1500 new features replacing the original 30 features. After the feature expansion step, we change the scale of data using sklearn's StandardScaler class.

## 5 Results

To show the accuracy of the method we split each sub-data-set into two train and test sets with the proportion of 75%-25%. In the table 4, we compared the approximation error with different loss functions. The results show that the proposed method can approximated the small house living spaces' with mean error of  $3.2m^2$  and  $4.1m^2$  for big houses. To show the accuracy of the model, we compute the following criteria:

$$accuracy = \frac{\text{no instances with true labels}}{\text{no total instances}} \quad (4)$$

in which a prediction label is true if it falls in the interval  $[0.85 * y_{true}, 1.15 * y_{true}]$ . The last column of table 4 reports the accuracy of the model.

Also, to show the error distribution, we drew box plot of residual errors for both small and big houses. The plots are drawn in the figure 4.

House Type	Mean Absolute Error	Mean Squared Error	Mean Epsilon Insensitive Error	Accuracy
Small	3.24	17.38	2.35	88.48%
Big	4.1	24.83	3.17	99.03%
Mean	3.67	21.11	2.76	93.76%
All	3.81	22.28	2.89	95.01%

Table 4: Approximation Error of living space feature




Figure 4: box plot for residual error of prediction for small and big houses

## 6 Conclusion

In this exercise we proposed a linear regression model equipped with a nonlinear feature expander, for predicting the living space area of different houses. Since our implementation of `LinearRegression` uses `scipy`'s `lstsq` function to solve a linear least squares problem, the time complexity of our model and the `sklearn`'s `LinearRegression` class are approximately equal.